

# Modern Text Analysis with Machine Learning: Week 2

Professor Anastasopoulos (ljanastas@uga.edu)

08/21/19

## For today

1. Example of discussion leader presentation.
2. Building and extracting corpora with APIs.
3. Pre-processing text data: tokenization, stemming, removing stop words.
4. Document-feature matrix and TF-IDF.

# Intro to natural language processing

- ▶ NLP is a term that is understood as a set of methods which map natural language units (words, sentences, paragraphs, etc) into a machine readable form.
- ▶ Once we can figure out how to represent language to a machine, we can then use statistical/mathematical tools to learn things about texts (without having to read them!).

# NLP Applications (Real World)

- ▶ SPAM Filters.
- ▶ Speech recognition (Siri).
- ▶ Machine translation.
- ▶ Information retrieval (search engines).
- ▶ Artificial intelligence.

# NLP Applications (Social Science)

- ▶ *Political sentiment and media bias* - Soroka, Stuart and Lori Young. 2012. "Affective News: The Automated Coding of Sentiment in Political Texts" *Political Communication* 29: 205-231.
- ▶ *Identification of politically relevant features of texts* - Monroe, Burt, Michael Colaresi, and Kevin Quinn. 2008. "Fightin' Words: Lexical Feature Selection and Evaluation for Identifying the Content of Political Conflict". *Political Analysis* 16(4)
- ▶ *Measuring expressed agendas in political texts* - Grimmer, J., 2010. A Bayesian hierarchical topic model for political texts: Measuring expressed agendas in Senate press releases. *Political Analysis*, pp.1-35.
- ▶ *And much much more...*

# NLP Basic Terminology

$$\textit{word} \subset \textit{document} \subset \textit{corpus}$$

- ▶ **document** - A collection of words, usually the unit of observation.
- ▶ **corpus** - A collection of documents or a single document.
- ▶ The **corpus** can be thought of as our entire dataset.
- ▶ The **document** can be thought of as an observation

## Example: Measuring political ideology from Tweets

Barberá, P., 2013. "Birds of the same feather tweet together, bayesian ideal point estimation using twitter data" *Political Analysis*

- ▶ Measuring the political ideology of Twitter users from tweets and patterns of following/followers.
- ▶ **document** - Each tweet is a document.
- ▶ **corpus** - All of the tweets that are analyzed are the corpus.

# Acquiring text data

1. *Online corpora* - There are many built in packages in **R** and in **Python** that you can load which have thousands of texts that you can access.
2. *Build your own corpora*
  - (a) Using an application program interface (API)
  - (b) Webscraping



## Online corpora

```
## Warning: package 'pacman' was built under R version 3.5.0
## Warning: unable to access index for repository http://datacube.wu.ac.at/bin/macosx/
## cannot open URL 'http://datacube.wu.ac.at/bin/macosx/
## installing the source package 'tm.corpus.Reuters21578'
```

```
library(tm)
install.packages("tm.corpus.Reuters21578",
                 repos = "http://datacube.wu.ac.at")
library(tm.corpus.Reuters21578)
data(Reuters21578)
```

- ▶ There are **thousands** of online corpora that are available.
- ▶ **R** is not really that good for accessing online corpora because they're not as easy to acquire.
- ▶ Here we are accessing the **Reuters** corpus which is a collection of 21,578 Reuters articles from the Reuters newswire in 1987.

## Reuters Corpus

- Let's use the **tm** package to see what these articles look like.

```
inspect(Reuters21578[1:2])
```

```
## <<VCorpus>>
## Metadata:  corpus specific: 0, document level (indexed): 0
## Content:   documents: 2
##
## [[1]]
## <<PlainTextDocument>>
## Metadata:  16
## Content:   chars: 2860
##
## [[2]]
## <<PlainTextDocument>>
## Metadata:  16
## Content:   chars: 438
```

# Reuters Corpus

- ▶ We can also read the articles to see how they're structured in **R**

```
Reuters21578[[1]]$content
```

```
## [1] "Showers continued throughout the week in\nthe Bahia
```

# Building your own corpora using APIs

- ▶ In general, ready-to-go corpora are not that useful.
- ▶ They contain old documents and are unlikely to have the info you want.
- ▶ For the most part, when you're doing text analysis, you'll have to build your own corpora.

## Building your own corpora using APIs

- ▶ The easiest way to do this is to extract data from one of the MILLIONS of APIs out there.
- ▶ APIs were originally designed to allow developers of apps to have constant streaming access to data.
- ▶ But they are a treasure trove of information of immesurable use to social scientists who know how to tap into them.

# Politics APIs

- ▶ **GovTrack.us API** - Get any information about Congress (bills, legislators, voting) over several Congresses.
- ▶ **OpenStates API** - Tons of information about state legislators, bills, voting, etc.
- ▶ **Opensecrets.org** - Money in politics and campaign finance database.
- ▶ **Twitter API** - Get streaming tweets from Twitter with user information etc.

## Extracting data from APIs using R

- ▶ Three packages are very useful for this purpose: **twitteR**, **httr**, and **jsonlite**.
- ▶ You first must establish *OAuth* credentials if you would like to access Tweets.
- ▶ You can find out how to do so here [Getting OAuth Credentials](#)

# Extracting Tweets using TwitterR

```
library(twitterR)
setup_twitter_oauth(
  consumer_key="PeVki23tnu1dNzYoG1pJS9kh0", # This is the
  consumer_secret="RksPKIMCK2s7UBpiFuIJzUzQwCkXKY6nVtChlD0x",
  access_token="18249358-ctVZdV7IpFxOcorUUdNqjeahb9Tb23yVoc",
  access_secret="46JNswM5I3Q6JwgDsUtyYKCMhs7yDY9jsrxslH55v"
)

## [1] "Using direct authentication"
```



# Extracting Tweets using TwitterR

Tweets are saved as a list in *R*

```
UGATweets = searchTwitter("@universityofga")[1:2]  
UGATweets[[1]]
```

```
## [1] "UGA_Innovation: .@uga_genetics @universityofga @UGA"
```

## Example 2: Tapping into APIs using “jsonlite”

## Example 2: Tapping into APIs using “jsonlite”

```
# We can retrieve the title and other information about the  
# I'm creating a data frame with the bill title, bill id, i  
# id and the bill sponsors gender
```

```
billtitles<-bills$objects$title  
billid<-bills$objects$id  
sponsorid<-bills$objects$sponsor$id  
sponsiridgender<-bills$objects$sponsor$gender
```

```
refugeebilldat<-  
  data.frame(billtitles,sponsiridgender)
```

```
#head(refugeebilldat)
```

## Pre-processing text data

- ▶ Now that we have the data, we need to prepare it for analysis.
- ▶ This involves a couple of steps which take **strings** and breaks them down into analyzeable units.

# Pre-processing text data steps

1. **Tokenization** - splits the document into tokens which can be words or n-grams (phrases).
2. **Formatting** - punctuation, numbers, case, spacing.
3. **Stop word removal** - removal of “stop words”
4. **Stemming** - removal of certain types of suffixes.

# Tokenization

- ▶ “**Bag of words**” model - most text analysis methods treat documents as a big bunch of words or terms.
- ▶ Order is generally not taken into account, just word and term frequencies.
- ▶ There are ways to parse documents into *ngrams* or *words* but we'll stick with words for now.

## Tokenization example: Tweets mentioning "@realDonaldTrump"

```
library(plyr)
library(quanteda) # This is a great text cleaning package

## Warning: package 'quanteda' was built under R version 3
## Package version: 1.4.0
## Parallel computing: 2 of 4 threads used.
## See https://quanteda.io for tutorials and examples.
##
## Attaching package: 'quanteda'
##
## The following objects are masked from 'package:tm':
##
##      as.DocumentTermMatrix, stopwords
##
## The following object is masked from 'package:utils':
##
```

Tokenization example: Tweets mentioning “@realDonaldTrump”

[illegible]



## Tokenization example: Tweets mentioning "@realDonaldTrump"

```
potustweets.vector.tokens[1:2]
```

```
## tokens from 2 documents.
```

```
## text1 :
```

```
## [1] "RT"
```

```
"JimKilbane"
```

```
## [3] "realDonaldTrump"
```

```
"RT_JimKilbane"
```

```
## [5] "JimKilbane_realDonaldTrump"
```

```
"RT_JimKilbane_realI
```

```
##
```

```
## text2 :
```

```
## [1] "RudyGiuliani"
```

```
## [2] "realDonaldTrump"
```

```
## [3] "They"
```

```
## [4] "swore"
```

```
## [5] "an"
```

```
## [6] "oath"
```

```
## [7] "to"
```

```
## [8] "the"
```

# Stop words

- ▶ Stop words are simply words that removed during text processing.
- ▶ They tend to be words that are very common “the”, “and”, “is” etc.
- ▶ These common words can cause problems for machine learning algorithms and search engines because they add noise.
- ▶ **BEWARE** Each package defines different lists of stop words and sometimes removal can decrease performance of supervised machine learning classifiers.

# Stemming

- ▶ In linguistics, stemming is the process of reducing words to their stems.
- ▶ “argue”, “argued”, “argues”, “arguing”, and “argus” reduce to the stem “argu”
- ▶ This is especially useful for unsupervised machine learning algorithms but may introduce issues in supervised machine learning.
- ▶ For example “cats” and “catty” would both be reduced to the term “cat”.

## Building the document-feature matrix

```
potustweets.dfm <- dfm(potustweets.vector.tokens, # Constr  
                        remove = stopwords("english"),  
                        stem = TRUE, remove_punct = TRUE)
```

```
potustweets.dfm
```

```
## Document-feature matrix of: 10 documents, 308 features
```

# Building the document-term matrix

```
potustweets.dfm[1:5, 1:5]
```

```
## Document-feature matrix of: 5 documents, 5 features (56
```

```
## 5 x 5 sparse Matrix of class "dfm"
```

```
##           features
```

```
## docs      rt jimkilban realdonaldtrump rt_jimkilban
```

```
## text1  1           1           1           1
```

```
## text2  0           0           1           0
```

```
## text3  1           0           1           0
```

```
## text4  1           0           1           0
```

```
## text5  0           0           1           0
```

```
##           features
```

```
## docs      jimkilban_realdonaldtrump
```

```
## text1           1
```

```
## text2           0
```

```
## text3           0
```

```
## text4           0
```

```
## text5           0
```

## Additional document-feature matrix options.

- ▶ There are a few other options that are useful for pre-processing a DFM in quanteda.
- ▶ *Dictionaries*: you can specify your own dictionary.
- ▶ *Remove*: you can remove specific words or patterns.
- ▶ *Automatic Sparsity Reduction* - An automated method for reducing the sparsity of a DFM. May help performance of ML algorithms.
- ▶ *Minimum/Max Number of Characters*: useful for topic modeling.
- ▶ *TF/IDF Weighting*: Multipurpose, useful primarily for *supervised learning*.

# Sparsity reduction

- ▶ Practically all DFMs are SPARSE (they are mostly 0s).
- ▶ Reducing the sparsity of a matrix can improve performance of machine learning algorithms.

## Automatic sparsity reduction code

```
potustweets.dfm.reduced = dfm_trim(potustweets.dfm, sparsiti  
potustweets.dfm.reduced
```

```
## Document-feature matrix of: 10 documents, 308 features
```



# TF/IDF Weighting

Term frequency:

$$tf_d = t_{d,i}$$

Inverse document frequency measures how much “information” a word contains.

$$idf_d = \log \frac{N}{n_i}$$